

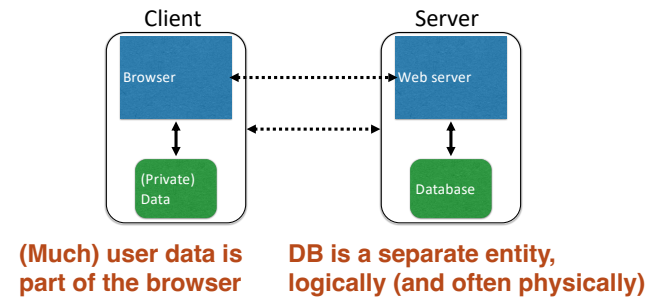
ECE/CS230

Computer Systems Security

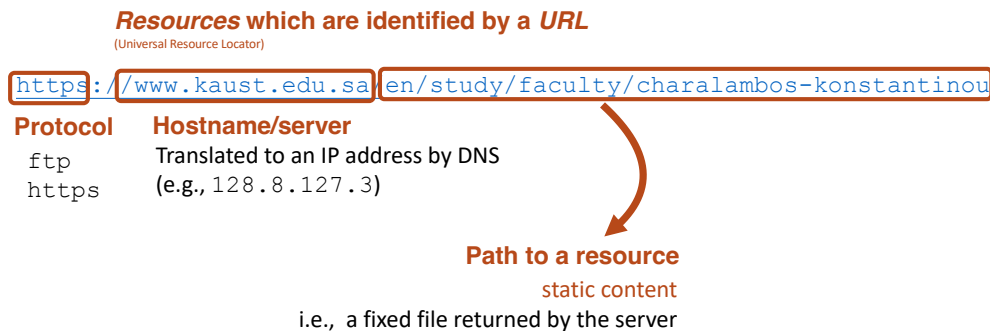
Charalambos (Harrys) Konstantinou
<https://sites.google.com/view/ececs230kaust>
Web

1

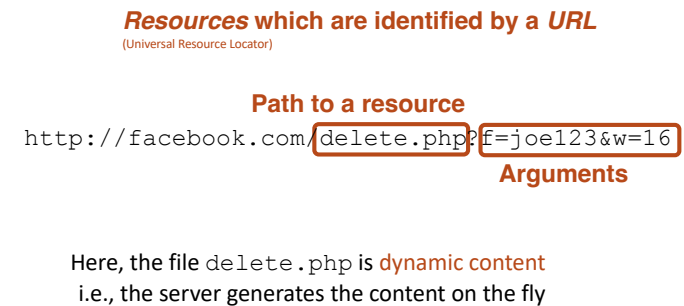
The web, basically




Interacting with web servers



Interacting with web servers



Basic structure of web traffic



- HyperText Transfer Protocol (**HTTP**)
 - An “application-layer” protocol for exchanging data

Basic structure of web traffic



- Requests contain:
 - The **URL** of the resource the client wishes to obtain
 - **Headers** describing what the browser can do
- Request types can be **GET** or **POST**
 - **GET**: all data is in the URL itself
 - **POST**: includes the data as separate fields

HTTP GET requests

```
GET /doc/test.html HTTP/1.1
Host: www.test101.com
Accept: image/gif, image/jpeg, */*
Accept-Language: en-us
Accept-Encoding: gzip, deflate
User-Agent: Mozilla/4.0
Content-Length: 35

bookId=12345&author=Tan+Ah+Teck
```

Annotations for the GET request example:

- Request Line: `GET /doc/test.html HTTP/1.1`
- Request Headers: `Host: www.test101.com`, `Accept: image/gif, image/jpeg, */*`, `Accept-Language: en-us`, `Accept-Encoding: gzip, deflate`, `User-Agent: Mozilla/4.0`, `Content-Length: 35`
- A blank line separates header & body
- Request Message Body: `bookId=12345&author=Tan+Ah+Teck`

HTTP POST requests

POST Request Example


```
POST /w/index.php?title=Wikipedia:Sandbox HTTP/1.1
Content-Type: application/x-www-form-urlencoded

wpStartTime=20080719022313&wpEditTime=20080719022100
...
```

Annotations for the POST request example:

- Blank line separates request headers and body
- Content type for data submitted via HTML form (multipart/form-data for file uploads)
- Request body: `wpStartTime=20080719022313&wpEditTime=20080719022100`
- Note: Most browsers have a query string length limit. Lowest known common denominator: IE7 `strlen(entire URL) <= 2,048 bytes`. This limit is not standardized. It applies to query strings, but **not request bodies**.

Basic structure of web traffic



- **Responses** contain:
 - **Status** code (<https://www.w3.org/Protocols/rfc2616/rfc2616-sec6.html>)
 - **Headers** describing what the server provides
 - **Data**
 - **Cookies** (much more on these later)
 - Represent *state* the server would like the browser to store

HTTP responses


	Status code
	HTTP/1.1 (200) OK
	Date: Sun, 18 Oct 2009 08:56:53 GMT
	Server: Apache/2.2.14 (Win32)
	Last-Modified: Sat, 20 Nov 2004 07:16:26 GMT
	ETag: "10000000565a5-2c-3e94b66c2e680"
	Accept-Ranges: bytes
	Content-Length: 44
	Connection: close
	Content-Type: text/html
	X-Pad: avoid browser bug
Header	
Data	<html><body><h1>It works!</h1></body></html>

Adding state to the web

HTTP is *stateless*

- The lifetime of an HTTP **session** is typically:
 - Client connects to the server
 - Client issues a request
 - Server responds
 - Client issues a request for something in the response
 - repeat
 - Client disconnects
- No direct way to ID a client from a previous session
 - So why don't you have to log in at every page load?

Maintaining State



- Web application maintains *ephemeral state*
- Server processing often produces intermediate results
- Send state to the client
- Client returns the state in subsequent responses

Two kinds of state: **hidden fields**, and **cookies**

Ex: Online ordering

socks.com/order.php socks.com/pay.php

Order

\$5.50

Order

Pay

The total cost is \$5.50. Confirm order?

Yes No

Separate page

Ex: Online ordering

What's presented to the user

```
pay.php
<html>
<head> <title>Pay</title> </head>
<body>
<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="5.50">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">
</body>
</html>
```

Ex: Online ordering

The corresponding backend processing

```
if (pay == yes && price != NULL)
{
    bill_creditcard(price);
    deliver_socks();
}
else
    display_transaction_cancelled_page();
```

Anyone see a problem here?

Ex: Online ordering

Client can change the value!

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" \value="0.01"
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Solution: *Capabilities*

- Server maintains *trusted* state
 - Server stores intermediate state
 - Send a pointer to that state (**capability**) to client
 - Client **references** the capability in next response
- Capabilities should be **hard to guess**
 - Large, random numbers
 - To prevent illegal access to the state

Using capabilities

Client can no longer change price

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="sid" value="781234">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Using capabilities


The corresponding backend processing

```
price = lookup(sid);
if (pay == yes && price != NULL)
{
    bill_creditcard(price);
    deliver_socks();
}
else
    display_transaction_cancelled_page();
```

But we don't want to use hidden fields all the time!

- Tedious to maintain on all the different pages
- Start all over on a return visit (after closing browser window)

Statefulness with Cookies



- Server maintains trusted state
 - Indexes it with a **cookie**
 - Sends cookie to the client, which stores it
 - Client returns it with subsequent queries to same server

Cookies

```
1 HTTP/1.0 200 OK
2 Content-type: text/html
3 Set-Cookie: yummy_cookie=choco
4 Set-Cookie: tasty_cookie=strawberry
5
6 [page content]
```

Now, with every new request to the server, the browser will send back all previously stored cookies to the server using the `Cookie` header.

```
1 GET /sample_page.html HTTP/1.1
2 Host: www.example.org
3 Cookie: yummy_cookie=choco; tasty_cookie=strawberry
```

Cookies are key-value pairs

Set-Cookie: **key=value; options;**

Headers

```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqca10cbiagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdnet-edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=zdnet.com
Set-Cookie: session-zdnet-production=590b97fprnge4ogese40vq11; path=/; domain=zdnet.com
Set-Cookie: user_agent=desktop
Set-Cookie: zdnet_ad_session=f
Set-Cookie: firstpg=0
Expires: Thu, 19 Nov 1981 08:52:00 GMT
Cache-Control: no-store, no-cache, must-revalidate, post-check=0, pre-check=0
Pragma: no-cache
X-UA-Compatible: IE=edge,chrome=1
Vary: Accept-Encoding
Content-Encoding: gzip
Content-Length: 18922
Keep-Alive: timeout=70, max=146
Connection: Keep-Alive
Content-Type: text/html; charset=UTF-8
```

Data

```
<html> ..... </html>
```

Cookies

Set-Cookie: **edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=zdnet.com**

Client

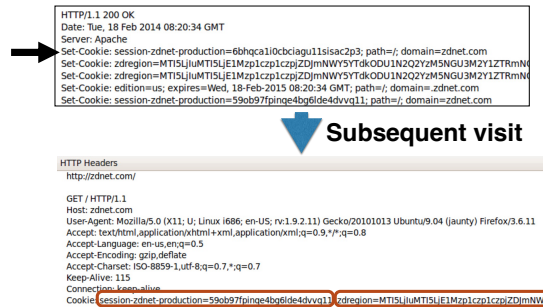
Browser

(Private) Data

Semantics

- Store “us” under the key “edition”
- This value was no good as of Feb 18, 2015
- This value should only be readable by any domain ending in `.zdnet.com`
- This should be available to any resource within a subdirectory of `/`
- Send the cookie with any future requests to `<domain>/<path>`

Requests with cookies



Why use cookies?

• Session identifier

- After a user has authenticated, subsequent actions provide a cookie
- So the user does not have to authenticate each time

• Personalization

- Let an anonymous user customize your site
- Store language choice, etc., in the cookie

Why use cookies?

• Tracking users

- Advertisers want to know your behavior
- Ideally build a profile *across different websites*
- Visit the Apple Store, then see iPad ads on Amazon?!
- How can site B know what you did on site A?

- Site A loads an ad from Site C
- Site C maintains cookie DB
- Site B also loads ad from Site C

- “Third-party cookie”
- Commonly used by large ad networks (doubleclick)

<http://live.wsj.com/video/how-advertisers-use-internet-cookies-to-track-you>

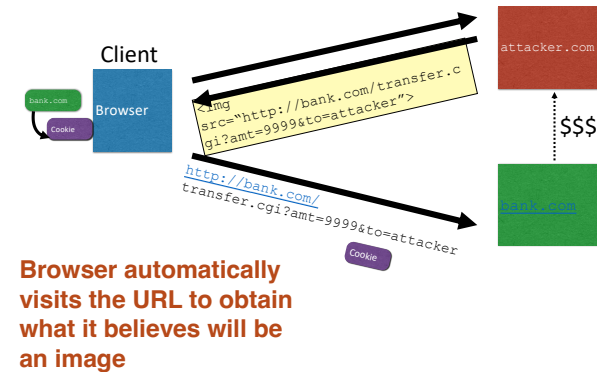
Cross-Site Request Forgery (CSRF)

URLs with side effects

`http://bank.com/transfer.cgi?amt=9999&to=attacker`

- GET requests often have **side effects on server state**
 - Even though they are not supposed to
- What happens if
 - the **user is logged in** with an active session cookie
 - a **request is issued for the above link?**
- How could you get a user to visit a link?

Exploiting URLs with side effects



Cross-Site Request Forgery

- **Target:** User who has an account on a vulnerable server
- **Attack goal:** Send requests to server *via the user's browser*
 - Look to the server like the user intended them
- **Attacker needs:** Ability to get the user to “click a link” crafted by the attacker that goes to the vulnerable site
- **Key tricks:**
 - Requests to the web server have predictable structure
 - Use e.g., `` to force victim to send it

Variation: Login CSRF

- Forge login request to honest site
 - Using **attacker's** username and password
- Victim visits the site under attacker's account
- What harm can this cause?

