

ECE/CS230

Computer Systems Security

Charalambos (Harrys) Konstantinou

<https://sites.google.com/view/ececs230kaust>

Program security

Before we start

- Assignment 1 grades are out
- Discussions are usual

Secure programs

- Why is it so hard to write secure programs?
- A simple answer:
 - Axiom (Murphy):
Programs have bugs
 - Corollary:
Security-relevant programs have security bugs

Outline

- Flaws, faults, and failures
- Unintentional security flaws
- Malicious code: Malware
- Other malicious code
- Nonmalicious flaws
- Controls against security flaws in programs

Outline

- Flaws, faults, and failures
- Unintentional security flaws
- Malicious code: Malware
- Other malicious code
- Nonmalicious flaws
- Controls against security flaws in programs

Flaws, faults, and failures

- A **flaw** is a problem with a program
- A **security flaw** is a problem that affects security in some way
 - Confidentiality, integrity, availability
- Flaws come in two types: **faults** and **failures**
- A fault is a mistake “behind the scenes”
 - An error in the code, data, specification, process, etc.
 - A fault is a **potential problem**

Flaws, faults, and failures

- A **failure** is when something actually goes wrong
 - You log in to the library's website, and it shows you someone else's account
 - "Goes wrong" means a deviation from the desired behavior, not necessarily from the specified behavior!
 - The specification itself may be wrong
- A **fault** is the programmer/specifier/inside view
- A **failure** is the user/outside view

Finding and fixing faults

- How do you find a fault?
 - If a user experiences a failure, you can try to work backwards to uncover the underlying fault
 - What about faults that haven't (yet) led to failures?
 - Intentionally try to **cause** failures, then proceed as above
 - Remember to think like an attacker!
- Once you find some faults, fix them
 - Usually by making small edits (**patches**) to the program
 - This is called “penetrate and patch”
 - “Patch Tuesday”* is a well-known example
https://en.wikipedia.org/wiki/Patch_Tuesday

Problems with patching

- Patching sometimes makes things **worse!**
- Why?
 - Pressure to patch a fault is often high, causing a narrow focus on the observed failure, instead of a broad look at what may be a more serious underlying problem
 - The fault may have caused other, unnoticed failures, and a partial fix may cause inconsistencies or other problems
 - The patch for this fault may introduce new faults, here or elsewhere!

Unexpected behavior (1/2)

- When a program's behavior is specified, the spec usually lists the things the program must do
 - The `ls` command must list the names of the files in the directory whose name is given on the command line, if the user has permissions to read that directory
- Most implementors wouldn't care if it did additional things as well
 - Sorting the list of filenames alphabetically before outputting them is fine

Unexpected behavior (2/2)

- But from a security / privacy point of view, extra behaviors could be bad!
 - After displaying the filenames, post the list to a public web site
 - After displaying the filenames, delete the files
- When implementing a security or privacy relevant program, you should consider “**and nothing else**” to be implicitly added to the spec
 - “should do” vs. “shouldn’t do”
 - How would you test for “shouldn’t do”?

Types of security flaws

- One way to divide up security flaws is by genesis (where they came from)
- Some flaws are intentional/inherent
 - **Malicious** flaws are intentionally inserted to attack systems, either in general, or certain systems in particular
 - If it's meant to attack some particular system, we call it a **targeted malicious flaw**
 - **Nonmalicious** (but intentional or inherent) flaws are often features that are meant to be in the system, and are correctly implemented, but nonetheless can cause a failure when used by an attacker
- Most security flaws are caused by unintentional program errors

Outline

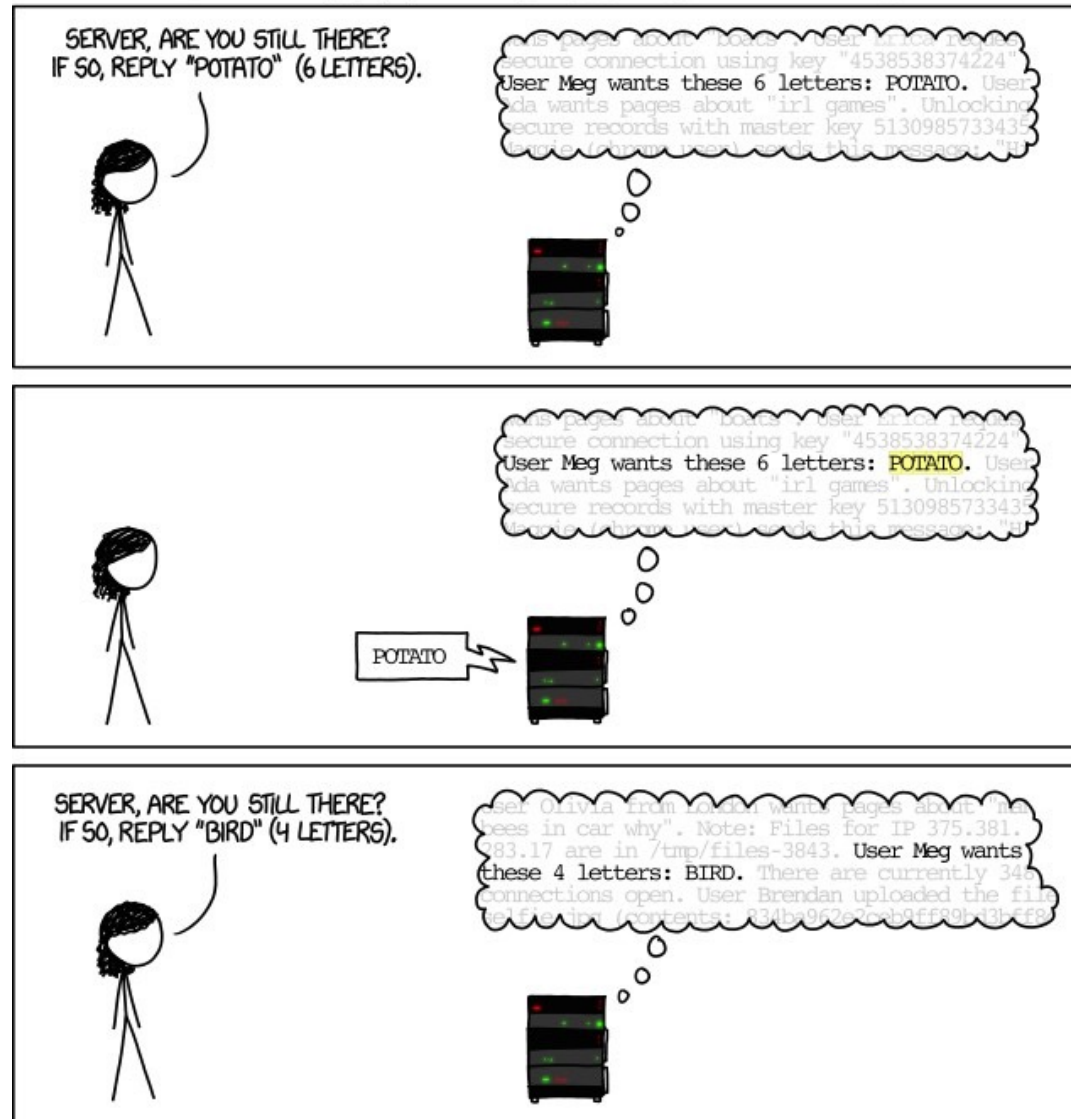
- Flaws, faults, and failures
- **Unintentional security flaws**
- Malicious code: Malware
- Other malicious code
- Nonmalicious flaws
- Controls against security flaws in programs

The Heartbleed Bug in OpenSSL (April 2014)

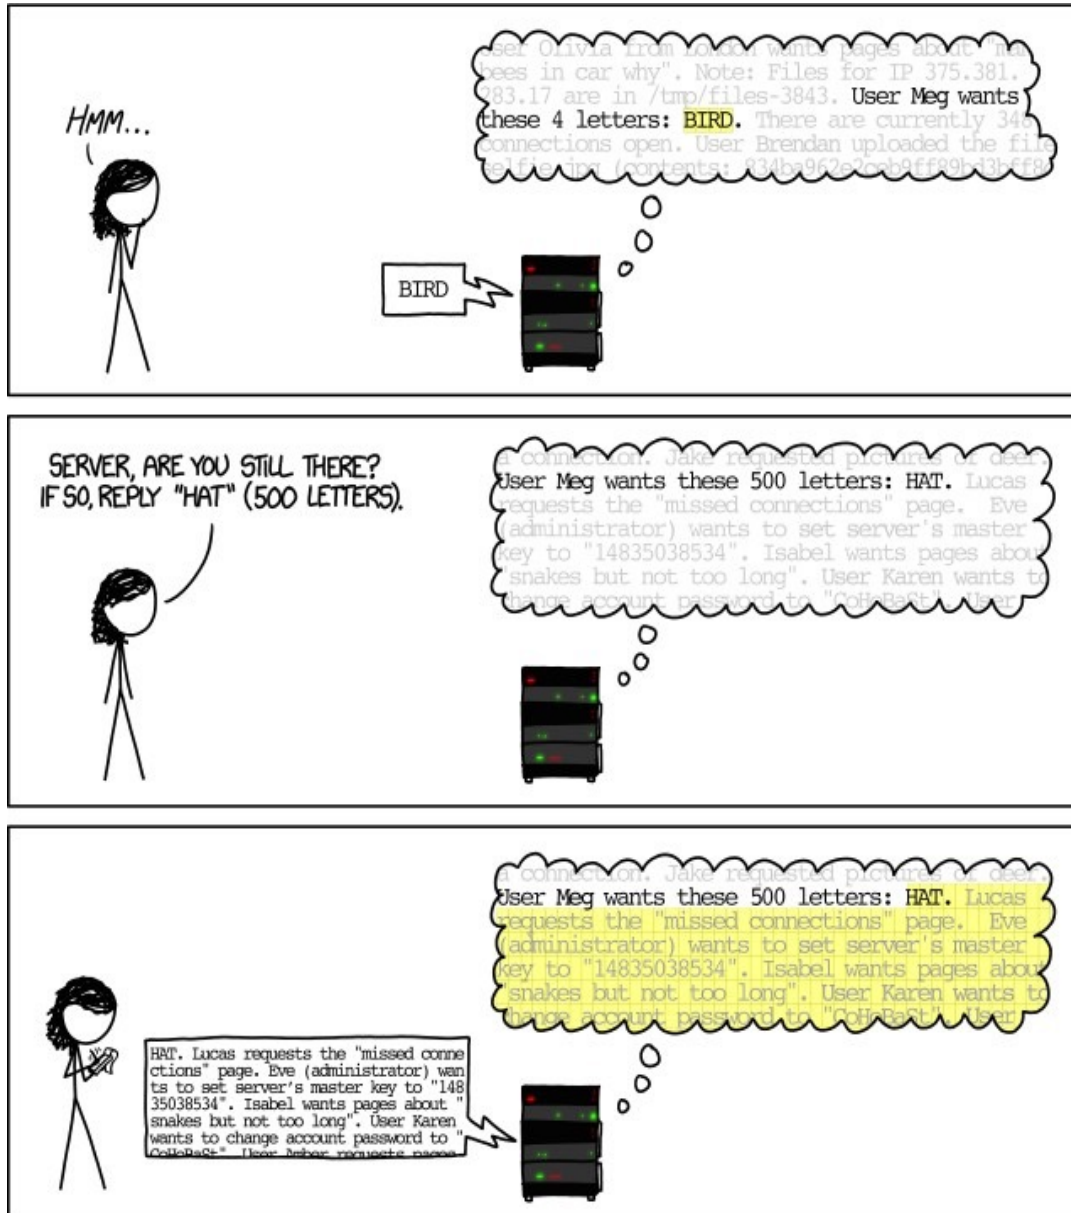
- The **TLS Heartbeat mechanism** is designed to keep SSL/TLS connections alive even when no data is being transmitted.
- Heartbeat messages sent by one peer contain random data and a payload length.
- The other peer is supposed to respond with a mirror of exactly the same data.

http://imgs.xkcd.com/comics/heartbleed_explanation.png

HOW THE HEARTBLEED BUG WORKS:



http://imgs.xkcd.com/comics/heartbleed_explanation.png



The Heartbleed Bug in OpenSSL (April 2014)

- There was a **missing bounds check** in the code.
- An attacker can request that a TLS server hand over a relatively large slice (up to 64KB) of its private memory space.
- This is the same memory space where OpenSSL also stores the server's private key material as well as TLS session keys.

Apple's SSL/TLS Bug (February 2014)

- The bug occurs in code that is used to check the validity of the server's signature on a key used in an SSL/TLS connection.
- This bug existed in certain versions of OSX 10.9 and iOS 6.1 and 7.0.
- An active attacker (a “man-in-the-middle”) could potentially exploit this flaw to get a user to accept a counterfeit key that was chosen by the attacker.

The Buggy Code

```
static OSStatus
SSLVerifySignedServerKeyExchange(SSLContext *ctx, bool isRsa, SSLBuffer signedParams,
                                uint8_t *signature, UInt16 signatureLen)
{
    OSStatus      err;
    ...

    if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
        goto fail;
    ...

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}
```

What's the Problem?

- There are two consecutive `goto fail` statements.
- The second `goto fail` statement is always executed if the first two checks succeed.
- In this case, the third check is bypassed and `0` is returned as the value of `err`.

Types of unintentional flaws

- Errors
- Integer overflows
- Buffer overflows
- Format string vulnerabilities
- Incomplete mediation
- TOCTTOU errors

Initialization Errors

- What happens when you use something that hasn't been initialized?
- This is a problem with languages like C/C++ and not Java
- This doesn't need to be a problem with C/C++, the programmer can make sure that they are initialized! (e.g. `int a = 0;` vs. `int a;`)
- This is a perfect example of trading performance over security/correctness.

Initialization Errors: Uninitialized variables

```
int count;  
while(count<100)  
{  
    cout<<count;  
    count++;  
}
```

Initialization Errors: variable to an uninitialized value

```
int a, b;  
int sum=a+b;  
cout<<"Enter two numbers to add: ";  
cin>>a;  
cin>>b;  
cout<<"The sum is: "<<sum;
```

When Run:

```
Enter two numbers to add: 1 3  
The sum is: -1393
```


Initialization Errors: fixed

```
int a, b;  
int sum;  
cout<<"Enter two numbers to add: ";  
cin>>a;  
cin>>b;  
sum=a+b;  
cout<<"The sum is: "<<sum;
```

Input Validation Error

- An Input Validation Error occurs when an input is NOT CHECKED to ensure it satisfies the assumptions (specifications)
- This is a general type of error (it overlaps with other types according to CWE)
- This is a very COMMON error

Input Validation Error – Example 1

So what happens to the following pseudo-code for the “myCat” program.

- The user inputs a filename/string. The filename is stored in strFilename.
- Execute the following command “cat [strFilename]” where [strFilename] is whatever the string is.

Input Validation Error – Example 1 Continued

- What does the command “myCat hello.txt” result in?
- According to pseudo-code
 - strFilename is now “hello.txt”
 - “cat hello.txt” is run
 - So the contents of hello.txt is printed out onto the screen.
- But remember this is about errors and vulnerabilities. See anything wrong?
- There wasn’t any “input validation”

Input Validation Error – Example 1 Continued

- In Unix commands are separated by ‘;’
- What happens when the command myCat “hello.txt;rm -rf /” (note that the thing in quotes is a single string).
- According to the pseudo-code:
 - strFilename is now “hello.txt;rm -rf /”
 - The command cat hello.txt;rm -rf / is now executed.
 - In actuality it is two commands (separated with ‘;’)
- Now add this with Unix (and Windows) systems not applying principle of least privilege then we have a formula for disaster

Input validation error - preconditions

- Lets take another example, this time with C/C++ code
//precondition: iaTemp is a non-null pointer to an array of ints. iLen is the number of elements in iaTemp.

```
void incArray (int* iaTemp, int iLen) {  
    for(int i=0;i<iLen;++i) {  
        iaTemp[i] = iaTemp[i] + 1;  
    }  
}
```

- ia = int array, i = int, str = string

Input validation error - preconditions

- Preconditions are good, right?
- Any one calling this function should abide by the preconditions right?
 - WRONG! Attackers are successful (most of the time) because they do things that are “unexpected” or beyond the specification
 - They simply don’t, and don’t have to, follow the rules!
 - Of course mistakes are made all the time as well

Input validation error - preconditions

- We need to VALIDATE our preconditions (which includes the assumptions/specifications made on the inputs)

- Add in

```
if (iaTemp == NULL)
{
    return;
}
```


Numeric Errors – Example

- A numeric error occurs when you misuse a “type” or “types”
- Example:

```
for (int i = 0; i < strlen(strTemp); ++i)
{
    printf(“%c”, strTemp[i]);
}
```
- What is the type of the THING that goes in between [and]? Similarly, what is the type of the THING that is returned by strlen()?
- int? WRONG
- size_t is the right type
 - Follow on question is: size_t = unsigned int?
 - NO! These definitions are library AND platform dependent. REMEMBER OUR ASSUMPTIONS?
 - What if it is unsigned int. Now what?

Numeric Errors – Example Continued

- The previous example is an example of a “type mismatch”
- Now what if `size_t = unsigned int`?
 - The comparison `i < strlen()` IS INCORRECT! Since `i` is “signed” it is only about half as many possible positive values as `size_t`!
- Due to this, the previous example can also be said to be an example of a “sign mismatch”

Integer Overflows - Example

- Lets say that int and unsigned int are 2 bytes long

```
unsigned int ui = 0xFFFF;  
int i = 0x7FFF;  
++ui;  
++i;
```

- So what is ui?
 - 0x0000//0 in decimal!!!
There was an overflow!!
- What is i?
 - 0x8000// -32768!!!
There is another overflow, and even worse the sign changed!!!

Integer Overflows - Example Continued

- So what is the big deal?
- It all depends on the code. We can call this a “fault” so if we are prepared for it, then everything is good. If we are not, then things could be bad.
- See some examples in the SAMATE database.
- Ohh this reminds me, what is the opposite of $>$? (that is the greater than operator)
 - If you answered $<$ you are very wrong and have just introduced a very common numerical error that leads to “off-by-one” errors among other things.
 - Right answer is \leq

Integer overflows – summary

- Machine integers can represent only a limited set of numbers, might not correspond to programmer's mental model
- Program assumes that integer is always positive, overflow will make (signed) integer wrap and become negative, which will violate assumption
 - Program casts large unsigned integer to signed integer
 - Result of a mathematical operation causes overflow
- Attacker can pass values to program that will trigger overflow

Buffer overflows

- Simply, a buffer overflow is an event where more stuff is being written into a buffer than the buffer is meant to hold.

Smashing The Stack For Fun And Profit

- This is a classic (read: somewhat dated) exposition of how buffer overflow attacks work.
- Upshot: if the attacker can write data past the end of an array on the stack, she can usually **overwrite things like the saved return address**. When the function returns, it will jump to any address of her choosing.
- Targets: programs on a local machine that run with setuid (superuser) privileges, or network daemons on a remote machine

Buffer overflows

- Simply, a buffer overflow is an event where more stuff is being written into a buffer than the buffer is meant to hold.
- Example:
 `char strTemp[10] = "1234567890";`
- What is the problem? Well lets do some counting.
- We defined the string strTemp to be 10 characters (bytes) long and then we assigned it to a 10 character string. So no problems right?
- No, because "1234567890" is a c-string, which means there is a NULL character at the end. So we are stuffing 11 bytes into a 10 byte buffer.
- This is known as an off-by-one error, but is indeed a buffer overflow

Buffer Overflows Continued

- So then what happens?
- Most of the time, for the example above, absolutely nothing!
- That is because compilers like to “pad” variables to certain boundaries (e.g. 4-byte boundaries).
- This means that the compiler would allocate 12 bytes for that 10 byte buffer. So there is in reality enough room for that extra NULL.
- But what happens when we put in MUCH more than what the buffer can hold? Lets find out.

Lok's Computer 1

- To facilitate this next talk, I will now introduce a new computer so we can readily see what a buffer overflow can result in, AND why these type of errors are so powerful (or devastating if you are on the receiving end)
- Lok's computer has lots of memory that is linear and is separated into bytes (byte-addressing).
 - Address starts at 0x00
 - Address ends at 0x7F (128 bytes total)
 - Data grows up
 - Organization: Data appears first in memory followed by 0xFF and then comes the instructions
 - (See example in a couple slides)

Lok's Computer 2

- The computer only supports one type: string
 - All strings end with NULL ('\0')
 - When defining a string, we can use the [] to specify an initial length.
 - All variable names are pointers (so they contain an address)
- The computer has a very simple instruction set:
 - read VAR. Reads a string from the user and inputs it into the memory space starting at address VAR
 - read has code "0xF1"
 - write VAR. Writes the string starting at memory address VAR to the screen
 - write has code "0xF2"
- The computer runs very simply as well: 1. Find first instruction, take in 2 bytes, first byte is instruction (interpret it) 2nd byte is whatever instruction is expecting. Once complete, move 2 more bytes down memory and process that instruction.

Lok's Computer 3

```
strHello[12] = "Hello World";
```

```
write strHello;
```

- So how does this look in memory?
 - Lets take first row is main memory
 - Things appearing in quotes are characters, '0' is the character for zero
 - Otherwise the values are in hex, 00 is the hex value zero like NULL
 - Second is the address

[illegible]

Lok's Computer – Example 4

[illegible]

- So what does the above program do?
Find the start of the instructions and interpret:
 - We find F2,00 as the first instruction, which is print out the string starting at address 00
 - Then the next instruction is F2,0A which is print out the string starting in address 0A
- So output on the screen is:
"Hi MynameisLok name"

Lok's Computer - Example 4 Cont.

- Notice that in the example, the instruction for write (F2) was used in a very strange way.
- F2,0A actually pointed to the MIDDLE of the string!!
- This is a usage that wasn't expected or at least didn't seem to have been designed into the computer, but it is indeed possible
- Remember the assumptions!
- Now lets continue with a READ this time.

Example 5

- Lets say we need a program that reads in a name that is less than 10 characters long and says hi.

```
strMessage = "Your Name?";  
strHello = "Hello ";  
strName[11]; //one extra character for NULL  
  
write strMessage;  
read strName;  
write strHello;  
write strName;
```

Example 5 Cont.

[illegible]

- Above is the before snapshot
- Below is after the user inputs “Lok”

[illegible]

Example 5 Cont.

- So! We are talking about buffer overflows right? Lets try to do an overflow and see what happens.
- What will memory look like if the user inputs “BufferOverflowsAreBad”
So what happens now? The last instruction we processed was F1,12 at address 20. But now

'Y'	'o'	'u'	'r'	' '	'n'	'a'	'm'	'e'	'?'	00	'H'	'e'	'l'	'l'	'o'	' '	00	'B'	'u'	'f'	'f'	'e'	'r'	'O'	'v'	'e'	'r'	'f'	'l'	'o'	'w'	's'	'A'	'r'	'e'	'B'	'a'	'd'	00							
00						06				0A					0F	10															1F	20														2D

- So what happens now? The last instruction we processed was F1,12 at address 20. So our computer here expects the next instruction to be at address 22.
 - The contents of 22 and 23 is now 'r' and 'e'. And 'r' is definitely neither F1 nor F2 so?
 - Fault! Computer burns up and we see smoke, and Lok is sad because he paid a lot of money for his computer

Example 5 Cont.

- You thought we were done huh? NOPE.
- We just noticed that we could completely violate the original assumption that the user name was going to be 10 characters long
- We made the computer crash and burn
- We should also notice that if we were slick, we could change our name to some specially crafted string so that instead of 'r' and 'e' we get a real instruction.

Example 5 Cont.

- So lets do it, what do you think will happen if the following string was inserted as the username? You must pretend that the stuff in [] is actually the character corresponding to the hex value within the []

“BufferOverflowsA”[F2][0E]

- Here is what we will get:
“Your name?”
“|o “
- Then the program ends, and everyone is happy

Example 5 Cont.

So what happened?

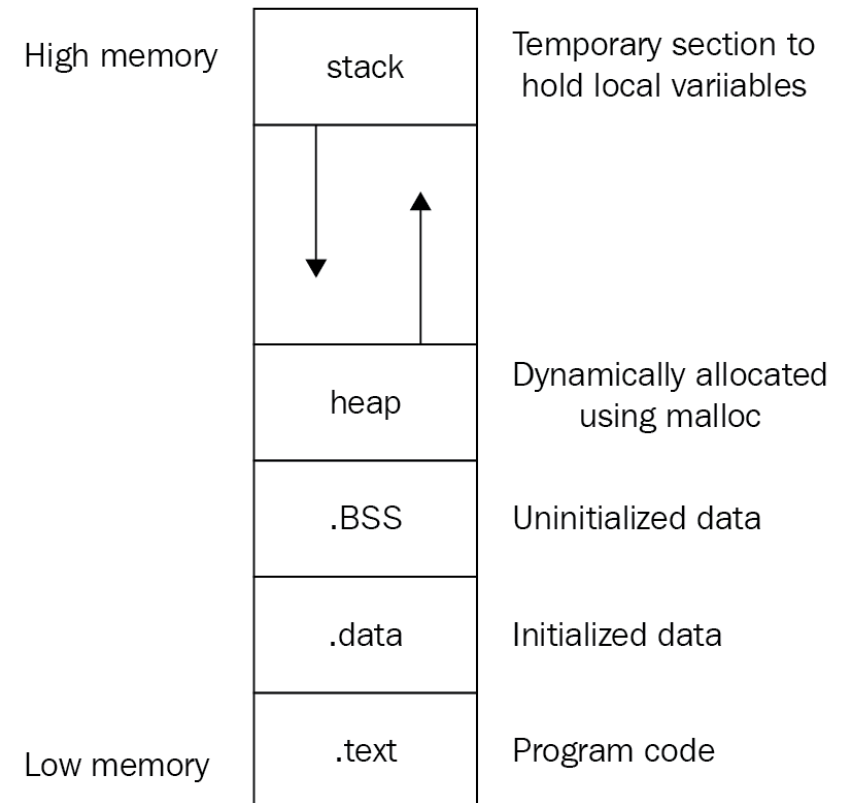
- We made an assumption that the user name was going to be a certain length.
- User was bad and inputted something really long.
- Turns out that it was SO long that it overwrote some instructions
- Lok's computer just kept going to the next expected instruction
- At first it was junk, but then we got smart and we were able to make Lok's computer do something the original program didn't do!
- This last point is the important part! Let's now move onto some real buffer overflow stuff

Buffer Overflows - Stack

- Now, we will focus on the stack, so stack overflows. The same concepts can be extended, but we don't talk about those

What is the stack? For x86

- Stack is on top of memory
- It grows down, but buffers go up (strange huh?)
- Stack not only contain local-variables, but also control data, “stack frames” for functions
- The last bullet is why stack overflows can be bad.



Intermission – Function Preamble / Postamble, Finale

- So we need to talk about the function preamble and finale (postamble? That is not a word I know) with respect to C and x86
- The preamble is a setup procedure whenever a function is called, the finale is the takedown

Preamble

When a function is called the following happens

- Function parameters are pushed onto the stack in reverse order
- The function is then called using the CALL instruction - The return address of the next instruction (the one after the function call) is pushed onto the stack. This is known as the return address
- The current stack frame base pointer (where the current frame starts) is pushed onto the stack

Postamble/Finale

When a function ends the following takes place

- The saved stack frame base pointer is restored to the appropriate register (EBP)
- The Instruction Pointer is set to the saved return address. This way the next instruction will be the one after the function was called.

Stack Frames (Simplified)

```
void f(int a)
{
    char strTemp[4];
}

void main()
{
    f(3);
}
```

- Left column is the address
- Right column is contents
- I made up the addresses, but they are labeled correctly
- RETURN ADD is the saved return address, i.e. the address of the next instruction after the CALL instruction that started the current function
 - The idea is so that when the current function ends, RETURN ADD is the next instruction to run
- EBP holds the base of the previous stack frame (Extended Base Pointer)

		main
0x7FFFFFFC	RETURN ADD	
0x7FFFFFF8	EBP	
0x7FFFFFF4	3	f
0x7FFFFEC	RETURN ADD	
0x7FFFFE8	EBP	
	strTemp[3]	
	strTemp[2]	
	strTemp[1]	
0x7FFFFE4	strTemp[0]	
0x7FFFFFF0		

Notice something?

- On Lok's computer we were able to overwrite the next instruction
- What about in the stack frame from the previous slide?
 - No instructions 😞
 - BUT! There is that return address.. Maybe if we can just overwrite that with another address
 - What other address?
 - Any address
 - What about an address on the stack, like strTemp?

Stack Overflow - Example 6

- So lets now just pretend that strTemp is user inputted.
- What does the user need to input in order to make that first return address point back to the start of strTemp?
- “abcdabcd”[0x7FFFFFFE4] is one possibility. There are plenty more.
 - That is interesting...
- What happens when this function ends?
 - The next instruction is now going to be 0x7FFFFFFE4.
 - But that is this “abcd” thing, so it will probably just crash. Just like in Lok’s computer.

		main
0x7FFFFFFC	RETURN ADD	
0x7FFFFFF8	EBP	
0x7FFFFFF4	3	f
0x7FFFFEC	RETURN ADD	
0x7FFFFE8	EBP	
	strTemp[3]	
	strTemp[2]	
	strTemp[1]	
0x7FFFFFFE4	strTemp[0]	
0x7FFFFFF0		

Example 6 Cont.

- So... if we were slick, then we will replace “abcd” with a valid instruction right?!
- Absolutely. Actually, if we were REALLY slick, we will replace it with a whole SERIES of instructions. Like instructions to open a “root shell”
 - A “root shell” is just a command prompt that has root privileges
 - Other possibilities are opening a port to listen to instructions
 - Anything else really

Shellcode

- The code to open a “shell” is known as – SHELLCODE
(See <http://shell-storm.org/shellcode/> for samples)
- Now that is interesting. But then there are a couple of difficult things.
 - First is how do we even know the address of strTemp in the first place?
 - Second is how do we know where the return address is, so I know how long my string is and also how to align things right?

Address of strTemp

- The address of where your “shellcode” will be is difficult to obtain
 - You might be able to get it through debugging the program
 - You can try trial and error
 - But in general, there are well known ways (we’ll talk about mitigation strategies, one of which is called address space layout randomization)
- Okay, so the start of the “shellcode” is difficult to pin down, but what if there was a way to make the start of the shellcode span a LARGE area?
 - There is. In x86, instruction 0x90 is the NOP instruction, which is exactly 1 byte (which is good) and does absolutely nothing.

Some stack

	0x7FFFFFFFC	0x7FFFFFF00
Address of Return Address →	0x7FFFFFF8	RETURN ADD
	0x7FFFFFF4	EBP
		buf[99]
		buf[98]
		...
Address of buffer →	0x7FFFE90	buf[0]

Stack with the buffer filled with 100 bytes of NOPs + SHELLCODE

