

ECE/CS230


Computer Systems Security

Charalambos (Harrys) Konstantinou

<https://sites.google.com/view/ececs230kaust>

Web

The web, basically



(Much) user data is part of the browser

DB is a separate entity, logically (and often physically)

Interacting with web servers

Resources which are identified by a *URL*

(Universal Resource Locator)

<https://www.kaust.edu.sa/en/study/faculty/charalambos-konstantinou>

Protocol

ftp
https

Hostname/server

Translated to an IP address by DNS
(e.g., 128.8.127.3)

Path to a resource

static content

i.e., a fixed file returned by the server

Interacting with web servers

Resources which are identified by a URL

(Universal Resource Locator)


Path to a resource

`http://facebook.com/delete.php?f=joe123&w=16`

Arguments


Here, the file `delete.php` is **dynamic content**
i.e., the server generates the content on the fly

Basic structure of web traffic



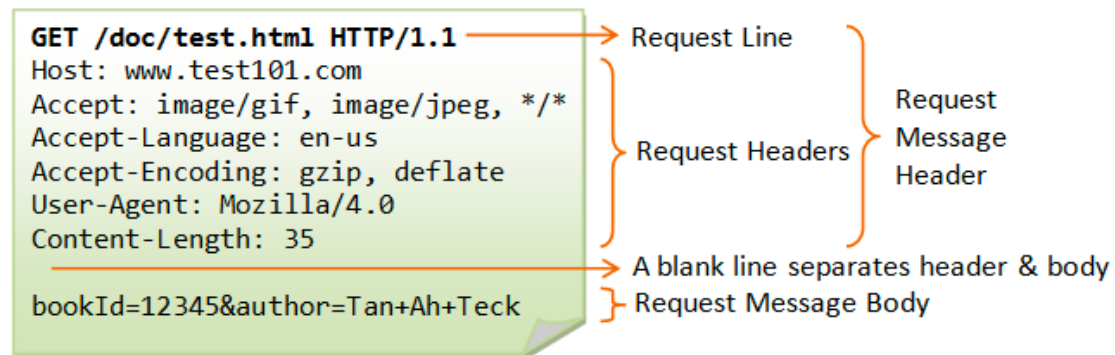
- HyperText Transfer Protocol (**HTTP**)
 - An “application-layer” protocol for exchanging data

Basic structure of web traffic



- Requests contain:
 - The **URL** of the resource the client wishes to obtain
 - **Headers** describing what the browser can do
- Request types can be **GET** or **POST**
 - **GET**: all data is in the URL itself
 - **POST**: includes the data as separate fields

HTTP GET requests



HTTP POST requests

POST Request Example

Blank line separates request headers and body

Content type for data submitted via HTML form (multipart/form-data for [file uploads](#))


```
POST /w/index.php?title=Wikipedia:Sandbox HTTP/1.1
Content-Type: application/x-www-form-urlencoded
wpStarttime=20080719022313&wpEdittime=20080719022100
...
```

Request body
... look familiar?

Note: Most browsers have a query string length limit. Lowest known common denominator: IE7
strlen(entire URL) <= 2,048 bytes.
This limit is not standardized. It applies to query strings, but **not request bodies**.



Basic structure of web traffic



- **Responses** contain:
 - **Status** code (<https://www.w3.org/Protocols/rfc2616/rfc2616-sec6.html>)
 - **Headers** describing what the server provides
 - **Data**
 - **Cookies** (much more on these later)
 - Represent *state* the server would like the browser to store

HTTP responses

Status code

Header

```
HTTP/1.1 200 OK
Date: Sun, 18 Oct 2009 08:56:53 GMT
Server: Apache/2.2.14 (Win32)
Last-Modified: Sat, 20 Nov 2004 07:16:26 GMT
ETag: "10000000565a5-2c-3e94b66c2e680"
Accept-Ranges: bytes
Content-Length: 44
Connection: close
Content-Type: text/html
X-Pad: avoid browser bug
```

Data


```
<html><body><h1>It works!</h1></body></html>
```

Adding state to the web

HTTP is *stateless*

- The lifetime of an HTTP **session** is typically:
 - Client connects to the server
 - Client issues a request
 - Server responds
 - Client issues a request for something in the response
 - repeat
 - Client disconnects
- No direct way to ID a client from a previous session
 - So why don't you have to log in at every page load?

Maintaining State





- **Web application maintains *ephemeral* state**
- Server processing often produces intermediate results
- Send state to the client
- Client returns the state in subsequent responses

Two kinds of state: **hidden fields**, and **cookies**

Ex: Online ordering

socks.com/order.php socks.com/pay.php



Separate page

Ex: Online ordering

What's presented to the user

pay.php

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="5.50">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Ex: Online ordering

The corresponding backend processing

```
if(pay == yes && price != NULL)
{
    bill_creditcard(price);
    deliver_socks();
}
else
    display_transaction_cancelled_page();
```

Anyone see a problem here?

Ex: Online ordering

Client can change the value!

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="0.01">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Solution: *Capabilities*

- Server maintains *trusted* state
 - Server stores intermediate state
 - Send a pointer to that state (*capability*) to client
 - Client **references** the capability in next response
- Capabilities should be **hard to guess**
 - Large, random numbers
 - To prevent illegal access to the state

Using capabilities

Client can no longer change price

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="sid" value="781234">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Using capabilities

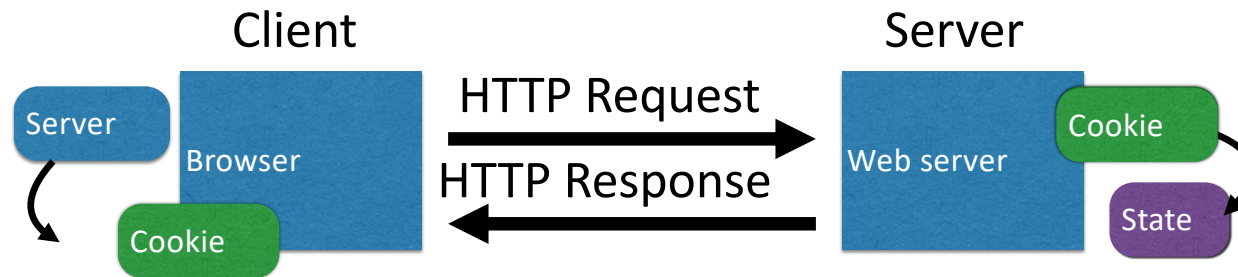
The corresponding backend processing

```
price = lookup(sid);  
if(pay == yes && price != NULL)  
{  
    bill_creditcard(price);  
    deliver_socks();  
}  
else  
    display_transaction_cancelled_page();
```

But we don't want to use hidden fields all the time!

- Tedious to maintain on all the different pages
- Start all over on a return visit (after closing browser window)

Statefulness with Cookies



- Server maintains trusted state
 - Indexes it with a **cookie**
 - Sends cookie to the client, which stores it
 - Client returns it with subsequent queries to same server

Cookies

```
1 | HTTP/1.0 200 OK
2 | Content-type: text/html
3 | Set-Cookie: yummy_cookie=choco
4 | Set-Cookie: tasty_cookie=strawberry
5 |
6 | [page content]
```

Now, with every new request to the server, the browser will send back all previously stored cookies to the server using the `Cookie` header.

```
1 | GET /sample_page.html HTTP/1.1
2 | Host: www.example.org
3 | Cookie: yummy_cookie=choco; tasty_cookie=strawberry
```

Cookies are key-value pairs

Set-Cookie: **key**=**value**; **options**;

Headers


Data

```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqca1i0cbciagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDlmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDlmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com
Set-Cookie: session-zdnet-production=590b97fpinqe4bg6ide4dvvq11; path=/; domain=zdnet.com
Set-Cookie: user_agent=desktop
Set-Cookie: zdnet_ad_session=f
Set-Cookie: firstpg=0
Expires: Thu, 19 Nov 1981 08:52:00 GMT
Cache-Control: no-store, no-cache, must-revalidate, post-check=0, pre-check=0
Pragma: no-cache
X-UA-Compatible: IE=edge,chrome=1
Vary: Accept-Encoding
Content-Encoding: gzip
Content-Length: 18922
Keep-Alive: timeout=70, max=146
Connection: Keep-Alive
Content-Type: text/html; charset=UTF-8

<html> ..... </html>
```

Cookies


Set-Cookie: `edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com`



Semantics

- Store "us" under the key "edition"
- This value was no good as of Feb 18, 2015
- This value should only be readable by any domain ending in `.zdnet.com`
- This should be available to any resource within a subdirectory of `/`
- Send the cookie with any future requests to `<domain>/<path>`

Requests with cookies



```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqca1i0cbciagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN4
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN4
Set-Cookie: edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com
Set-Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11; path=/; domain=zdnet.com
```



Subsequent visit

HTTP Headers

http://zdnet.com/

GET / HTTP/1.1

Host: zdnet.com

User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.9.2.11) Gecko/20101013 Ubuntu/9.04 (jaunty) Firefox/3.6.11

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8

Accept-Language: en-us,en;q=0.5

Accept-Encoding: gzip,deflate

Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7

Keep-Alive: 115

Connection: keep-alive

Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11; zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN4

Why use cookies?

- **Session identifier**

- After a user has authenticated, subsequent actions provide a cookie
- So the user does not have to authenticate each time

- **Personalization**

- Let an anonymous user customize your site
- Store language choice, etc., in the cookie

Why use cookies?

- **Tracking users**

- Advertisers want to know your behavior
- Ideally build a profile *across different websites*
- Visit the Apple Store, then see iPad ads on Amazon?!
- How can site B know what you did on site A?

- Site A loads an ad from Site C
- Site C maintains cookie DB
- Site B also loads ad from Site C

- **“Third-party cookie”**
- **Commonly used by large ad networks (doubleclick)**

Cross-Site Request Forgery (CSRF)

URLs with side effects

<http://bank.com/transfer.cgi?amt=9999&to=attacker>

- GET requests often have **side effects on server state**
 - Even though they are not supposed to
- What happens if
 - the **user is logged in** with an active session cookie
 - a **request is issued for the above link?**
- How could you get a user to visit a link?

Exploiting URLs with side effects



Browser automatically visits the URL to obtain what it believes will be an image

Cross-Site Request Forgery

- **Target:** User who has an account on a vulnerable server
- **Attack goal:** Send requests to server *via the user's browser*
 - Look to the server like the user intended them
- **Attacker needs:** Ability to get the user to “click a link” crafted by the attacker that goes to the vulnerable site
- **Key tricks:**
 - Requests to the web server have predictable structure
 - Use e.g., `` to force victim to send it

Variation: Login CSRF

- Forge login request to honest site
 - Using **attacker's** username and password
- Victim visits the site under attacker's account
- What harm can this cause?

